




A Research on the Method of Automatic Curriculum Learning for Reinforcement Learning

Weizhe Chen

Motivation

- Reinforcement Learning has succeeded in multiple domains

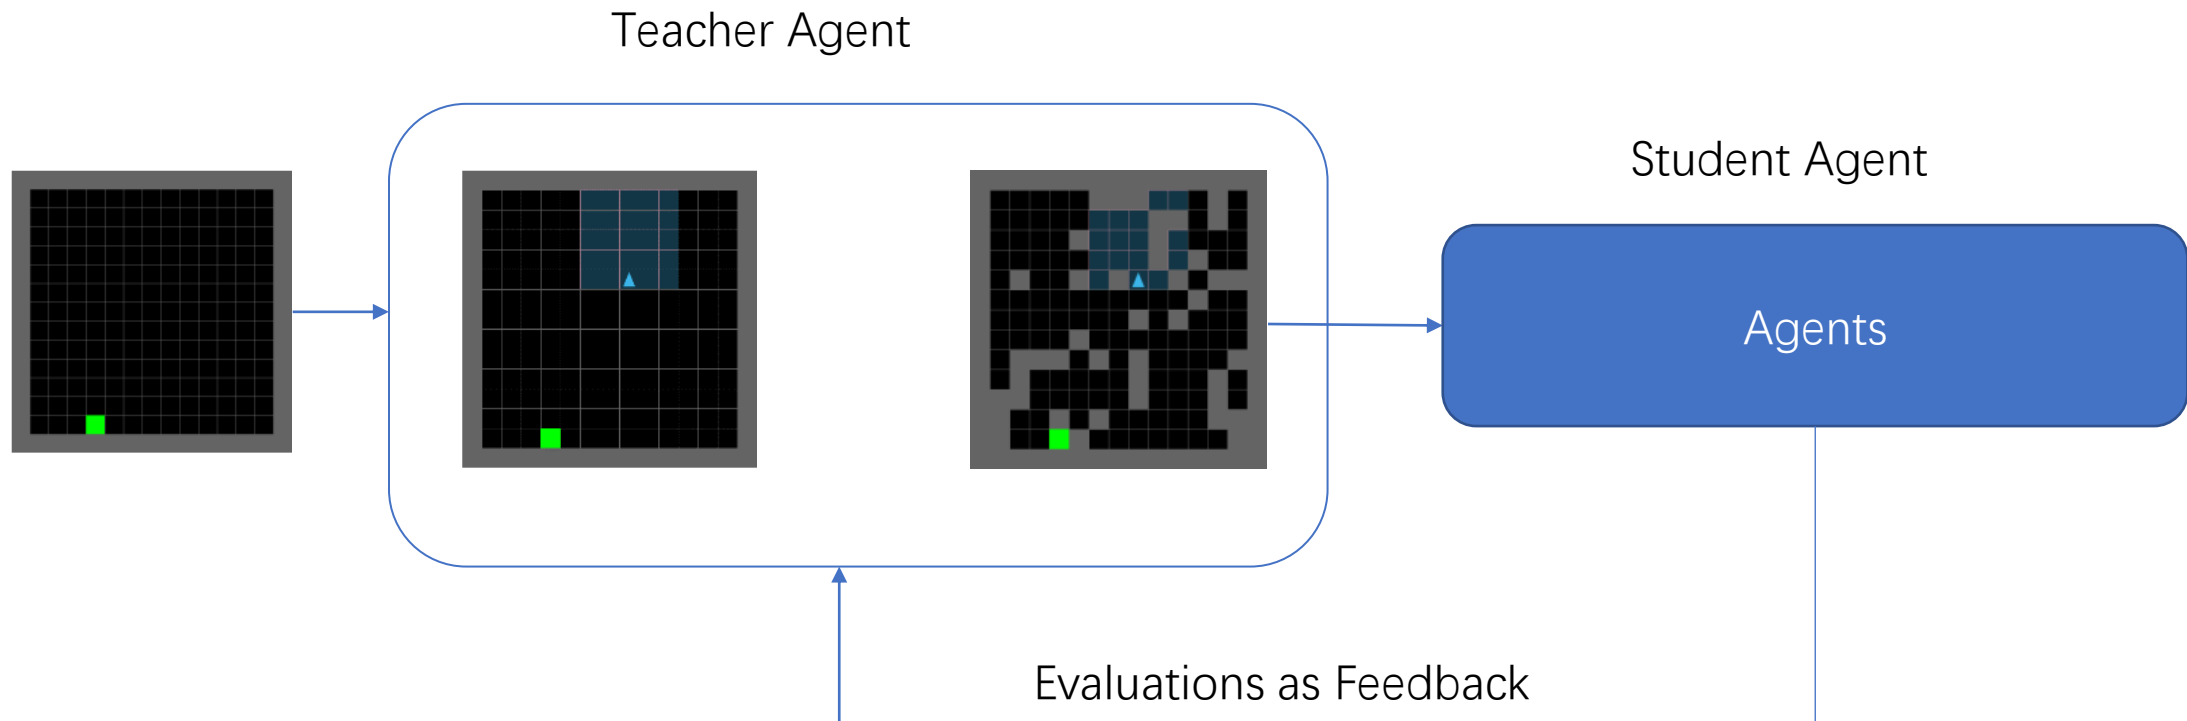


- Curriculum learning is one of the most important auxiliary tool behind
 - Turn a hard enough problem into a series of sub-task and learn them sequentially


The computing power is huge

- Dota2: 50k CPU, 500 GPU
- 2 vs 2 football: 4096+128 CPU, 16 TPU, 50 days

Unsupervised Environment Design



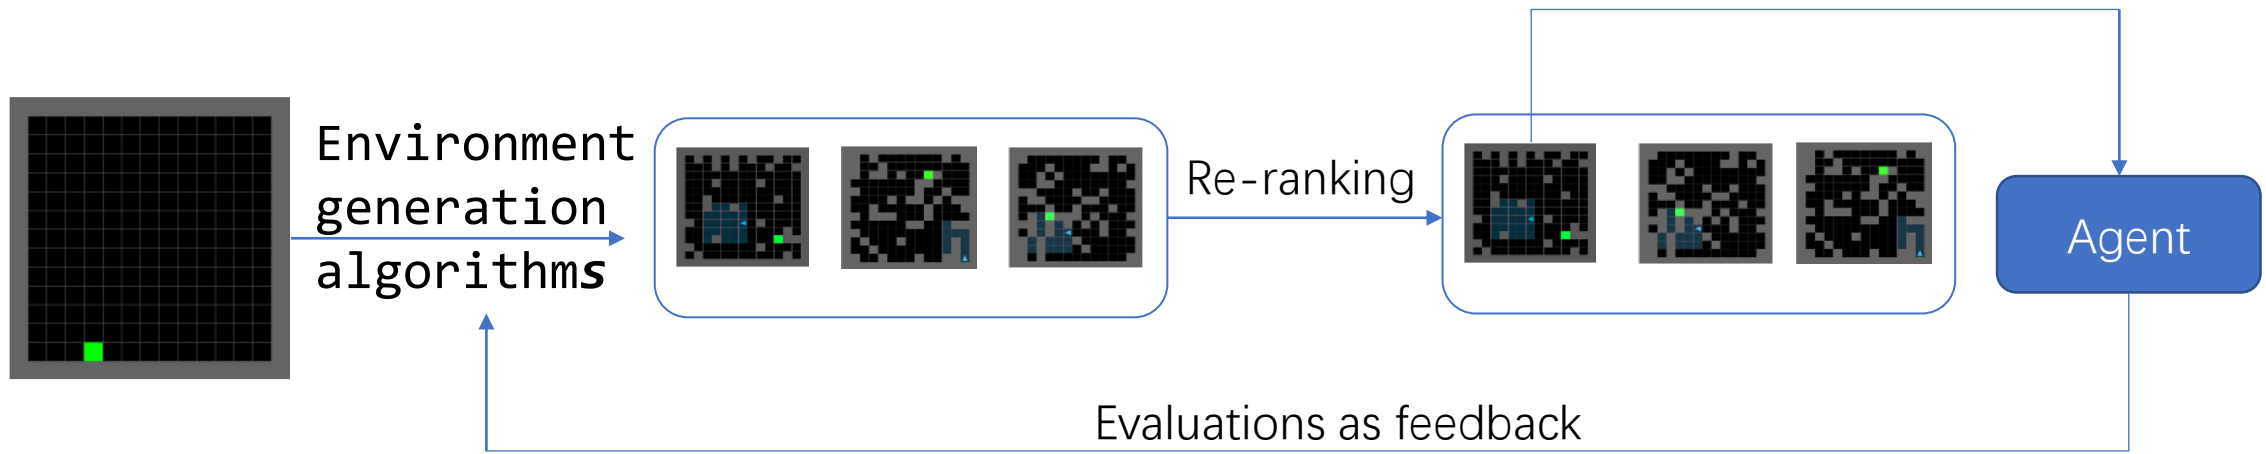
PAIRED



PAIRED

- Using the additional agent to calculate the regret value.
 - $U = \text{reward}(\text{Antagonist}) - \text{reward}(\text{Protagonist})$
 - $\text{Reward}'(\text{Antagonist}) = U$
 - $\text{Reward}'(\text{Protagonist}) = -U$
 - $\text{Reward}(\text{Environment}) = U$
- Making the environment to be easier for the agent to learn since the difference of a good policy and the reward of a bad policy is now larger

Problem Definition



Problem Definition

- A special Re-ranking Problem:
 - Re-ranking: one step in the recommendation system to rank the items so that the recommendation can consider more than just the predicted score but also diversity
 - Environment: Item
 - Agent: User

Markov Metric

- Defined as the difference between states
- Normal Requirement: $|V(s) - V(t)| \leq d(s, t)$
- well-setup metrics are hard to calculate!
- In this paper, we define $d(s, t) = |V(s) - V(t)|$
 - Can use the trained value function

SOPED

- From intuition, it is easier for agent to learn if the current environment is not changing too much from the previous one
- We re-rank the potential environments by the last environment we used in training
- Similarity Ordered Population-based Environment Design (SOPED)

```

1 Randomly initialize  $\pi^P, \pi^A, \pi^\Lambda$ ;
2 while Not converged do
3     Generate the environment parameters using the teacher agents population:  $\theta_i \sim \Theta_i \sim \pi_i^\Lambda$ ;
4     Calculate the initial observation  $o_i = I^{\theta_i}(s_i)$ ;
5     Calculate the distance with the latest trained environment  $M_{\theta^{r-1}}$ , and choosing the corresponding  $\theta_i$ 
        as the environment parameters used in this round of training  $\theta^r$ ;
6     Create the training environment  $M_\theta$  according to  $\theta_r$ ;
7     Sample training trajectories  $\tau^A$  and  $\tau^P$  in  $M_\theta$ , and calculate  $U(\pi^A)$  和  $U(\pi^P)$ ;
8     Calculate  $Regret = U(\pi^A) - U(\pi^P)$ ;
9     Train policy  $\pi^P$  with reward  $-Regret$ ;
10    Train policies  $\pi^A$  and  $\pi^\Lambda$  with reward  $Regret$ ;
11 end
12 return  $\pi^P, \pi^A, \pi^\Lambda$ 

```

SDOPED


- On the other hand, SOPED will more likely make the high-level generator to stuck in a local optimal rather than a global optimal
- We need to consider diversity during our training
- Determinantal Point Process (DPP) is used
 - Given the similarity of each pair of items and the score of the items, DPP can re-rank the sequence so that it both consider the score and the diversity
 - Similarity $s = e^{-0.5d}$, where d is the distance
- Similarity and Diversity Ordered Population-based Environment Design, SDOPED

SDOPED

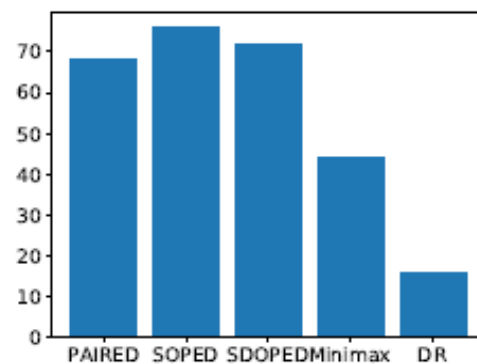

```
1 Randomly initialize  $\pi^P, \pi^A, \pi^\Lambda$ ;  
2 while Not converged do  
3   Generate the environment parameters using the teacher agents population:  $\theta_i \sim \Theta_i \sim \pi_i^\Lambda$ ;  
4   Calculate the initial observation  $o_i = I^{\theta_i}(s_i)$ ;  
5   Calculate the distance with the latest trained environment  $M_{\theta^{r-1}}$  as the score of the item;  
6   Calculating the similarity of each pair of environments as S;  
7   Calculate the Kernel Matrix  $L = \text{Diag}(r) \dot{S} \dot{\text{Diag}}(r) \Theta_{small} = \text{DPP}(L)$ ;  
8   Uniformly sample  $\theta \in \Theta_{small}$ ;  
9   Create actual environment  $M_\theta$  according to  $\theta_r$ ;  
10  Sample training trajectories  $\tau^A$  and  $\tau^P$  in  $M_\theta$ , and calculate  $U(\pi^A)$  和  $U(\pi^P)$ ;  
11  Calculate  $\text{Regret} = U(\pi^A) - U(\pi^P)$ ;  
12  Train policy  $\pi^P$  with reward  $-\text{Regret}$ ;  
13  Train policies  $\pi^A$  and  $\pi^\Lambda$  with reward  $\text{Regret}$ ;  
14 end  
15 return  $\pi^P, \pi^A, \pi^\Lambda$ 
```

Experiment

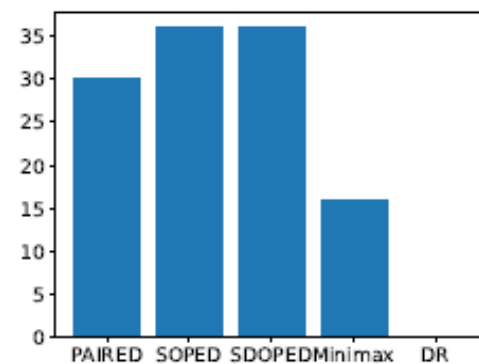

- Maze Design
 - Solved Path Length:





Experiment



(a). 16 Room




(b). Labyrinth



(c). Maze

Experiment

- Value loss:



Takeaway

- Using environment design, we can improve the training efficiency of the agents, and make the final result more robust
- The value predictor of the high-level environment generator can be easily trained, but the environment generator is very hard to train
- Even if the generator seems to have converge, it still could suddenly completely change its output